



# Machine Learning

Abdelhak Mahmoudi  
[abdelhak.mahmoudi@um5.ac.ma](mailto:abdelhak.mahmoudi@um5.ac.ma)

2021

# Content

## 1. The Big Picture

## 2. Supervised Learning

- Linear Regression, Logistic Regression, Support Vector Machines, Trees, Random Forests, Boosting, Artificial Neural Networks

## 3. **Unsupervised Learning**


- Principal Component Analysis, K-means, Mean Shift

# Unsupervised Learning

- **Dimensionality Reduction**
  - **Principal Component Analysis (PCA)**
- Clustering
  - K-Means
  - Mean-Shift



# Dimensionality Reduction

- Curse of dimensionality ( $n \gg m$ )
  - Data are at risk of being **very sparse** in high dimensional space
  - High risk of **overfitting**




# Dimensionality Reduction

- Transforms feature space from  $n$  to  $k$  ( $k < n$ )
  - Some **features** are probably **corelated** (dependent)
  - Some **features** are almost **constant**
  - **Transform but preserve** the maximum of **variance**

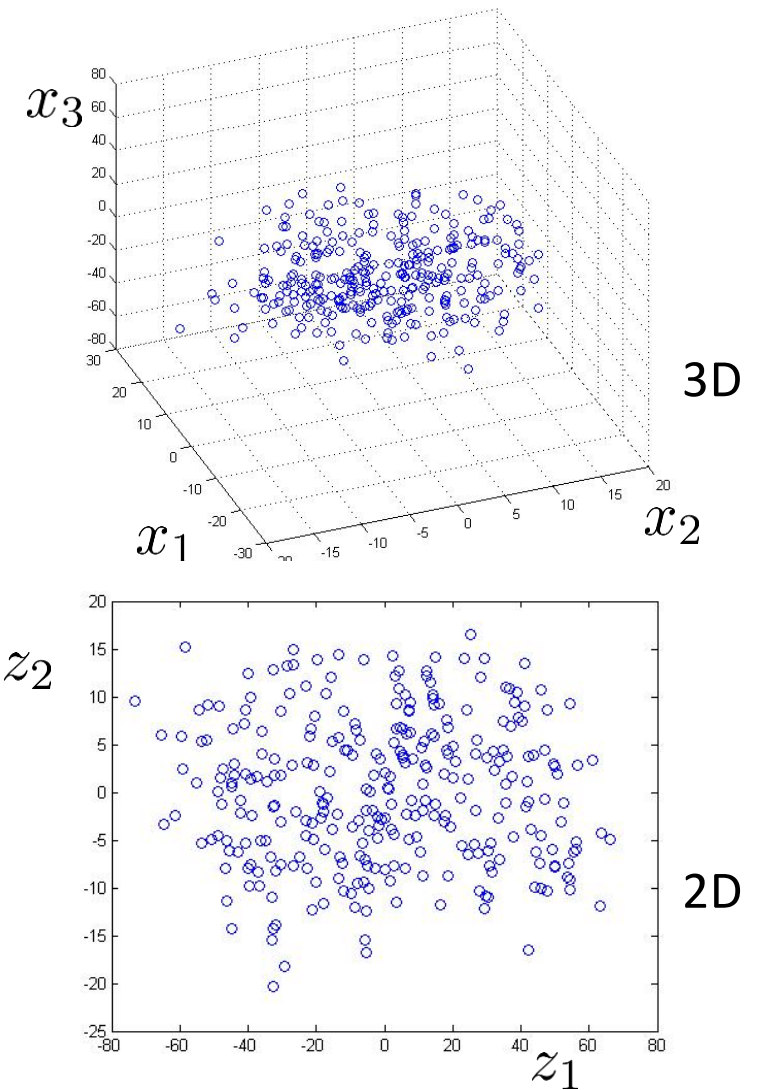



# Dimensionality Reduction


- Often
  - Not necessarily lead to better performance
  - Not the better way to address overfitting !
- Always
  - Speed up training
  - Allow data compression
  - Allow data exploration
  - Allow data visualization (DataViz)




# Principal Component Analysis



# Principal Component Analysis



**Maximum** loss  
Less variance



**Minimum** loss  
More variance



# Principal Component Analysis

- Singular Value Decomposition (SVD) (very costly)
  - Parallelization: Incremental PCA (fast), Randomized PCA (faster)
- PCA assumes that the dataset is centered around the origin
- How many dimensions to preserve?
  - Reduce dimensions that add up to a sufficiently large portion of the explained variance (e.g., 99%)
- Kernel PCA (kPCA): use the kernel trick like SVM
- In practice, use kPCA to transform the feature space, then perform classification or regression or clustering.

# Principal Component Analysis

- Hyper-Parameters Tuning
  - $d$ : polynomial Kernel
  - $\gamma$ : RBF kernel
  - $k$ : Number of retained principal components
  - Etc.

# Unsupervised Learning


- Dimensionality Reduction
  - Principal Component Analysis (PCA)
- **Clustering**
  - **K-Means**
  - **Mean-Shift**

# K-Means

- Pick a number of clusters **k**
- Initialize **centroids** randomly
- Problem of **local optima**
  - Run K-means a lot of times
- **Sensible** to **initial** conditions
- Have to **specify k** !

Repeat until convergence:

**Assign** each example to the cluster of the **nearest** centroid  
Compute the **mean** in each cluster  
Put the mean as the **new centroid**




# Mean Shift

- Chose a **radius  $r$  of the clusters**
- Initialize **centroids** at each example
- **No need** to specify the number of clusters

For each example:

Repeat until convergence:

Compute the **mean** in its cluster with radius  $r$   
**Shift** the cluster to the new mean centroid



# Other Clustering methods

- Expectation Maximization (EM)
- Hierarchical Clustering
- Affinity Propagation (AP)
- Etc.